

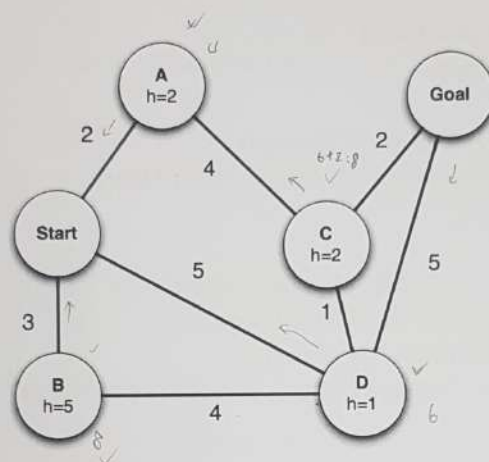
Due: 03/31/2022 23:59

Foundations of Artificial Intelligence: Homework 1

Instructor: Shang-Tse Chen & Hsuan-Tien Lin

Problem 1

(10 points)



Write down the order of (state expansion) and the final path returned by each of the graph search (as oppose to tree search) algorithms below. You can assume ties are resolved alphabetically.

- Depth-first search.
expansion: Start \rightarrow A \rightarrow C \rightarrow D \rightarrow B (back to D) \rightarrow Goal
final path: Start \rightarrow A \rightarrow C \rightarrow D \rightarrow Goal
- Breadth-first search.
expansion: Expand Start (discover A, B, D) \rightarrow Expand A (discover C) \rightarrow Expand B \rightarrow Expand D (discover Goal)
final path: Start \rightarrow D \rightarrow Goal
- Uniform cost search.
expansion: Expand Start (discover A, B, D) \rightarrow Expand A (discover C) \rightarrow Expand B \rightarrow Expand D (discover Goal)
 \rightarrow Expand C (update Goal) \rightarrow Goal
final path: Start \rightarrow A \rightarrow C \rightarrow Goal
- Greedy search with the heuristic h shown on the graph.
expansion: Start \rightarrow D \rightarrow Goal
final path: Start \rightarrow D \rightarrow Goal
- A* search with the same heuristic.
expansion: Expand Start (discover A, B, D) \rightarrow Expand A (discover C) \rightarrow Expand D (discover Goal)
 \rightarrow Expand B \rightarrow Expand C (update Goal) \rightarrow Goal

Problem 2

(10 points)

final path: Start \rightarrow A \rightarrow C \rightarrow Goal

Suppose that the heuristic (overestimates) the shortest path from any state to the goal by a factor of at most ϵ , where $\epsilon > 1$. Prove that the cost of the path found by A* tree search is at most ϵ times the cost of the optimal path. Suppose T_{true} is the cost announced by A*, S_{true} is the optimal cost,

T_{eval} and S_{eval} are their evaluation costs.

We have $\begin{cases} T_{eval} = T_{true} \text{ (termination step, the heuristic} = 0) \\ T_{eval} \leq S_{eval} \text{ (T terminates first)} \\ S_{eval} \leq \epsilon S_{true} \text{ (heuristic is } \epsilon\text{-bounded)} \end{cases}$

$$\Rightarrow T_{true} = T_{eval} \leq S_{eval} \leq \epsilon S_{true}$$

Hence, the cost found by A* is ϵ -bounded.

(10 points)

Problem 3

```

function A* GRAPH SEARCH(problem)
  fringe ← an empty priority queue
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  closed ← an empty set
  ADD INITIAL-STATE[problem] to closed
  loop
    if fringe is empty then
      return failure
    end if
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then
      return node
    end if
    for successor in GETSUCCESSORS(problem, STATE[node]) do
      if successor not in closed then
        ADD successor to closed
        fringe ← INSERT(MAKE-SUCCESSOR-NODE(successor, node), fringe)
      end if
    end for
  end loop
end function

```

Chances are a successor ^{could} have a better solution with a later expanded predecessor. But the implementation above simply ignore the successor if it is closed, which may lead to non-optimal solution.

The implementation of the A* graph search algorithm above is incorrect. Briefly explain the bug in this implementation and justify your answer.

Take Problem 1 as an example, the above implementation will return Start → D → Goal as the answer, which is not an optimal solution. (10 points)

Problem 4

You are scheduling for 6 classes taught by 3 instructors. Of course, each instructor can only teach one class at a time.

The classes are:

- Class 1 - Intro to Programming: 8:00-9:00am
- Class 2 - Intro to Artificial Intelligence: 8:30-9:30am
- Class 3 - Natural Language Processing: 9:00-10:00am
- Class 4 - Computer Vision: 9:00-10:00am
- Class 5 - Machine Learning: 10:30-11:30am

The instructors are:

- Instructor A - Can teach Classes 1, 2, and 5.
- Instructor B - Can teach Classes 3, 4, and 5.
- Instructor C - Can teach Classes 1, 3, and 4.

(1) Formulate this problem as a CSP. Describe the variables, domains and constraints.

Variables: Class 1 ~ Class 5

Domains: Class 1: {Instructor A, Instructor C}

Class 2: {Instructor A}

Class 3: {Instructor B, Instructor C}

Class 4: {Instructor B, Instructor C}

Class 5: {Instructor A, Instructor B}


Constraints: Class 1 \neq Class 2

Class 2 \neq Class 3


Class 2 \neq Class 4

Class 3 \neq Class 4

(2) Draw the constraint graph associated with your CSP.



(3) Show the domains of the variables after running arc-consistency on this initial graph (after having already enforced any unary constraints).



(4) Give one solution to this CSP.

Class 1: Instructor C

Class 2: Instructor A

Class 3: Instructor B

Class 4: Instructor C

Class 5: Instructor A

(5) Your CSP should look nearly tree-structured. Briefly explain (one sentence or less) why we might prefer to solve tree-structures CSPs.

For normal CSP, we may solve it in $O(d^n)$ time
 For tree-structured CSP, we may solve it in $O(nd^2)$ time